



# Northeastern University

## R Practice 4 ALY6015

R practice week 4 - Module 4 Assignment — Regularization

by Mohammad Hossein Movahedi

[movahed.m@northeastern.edu](mailto:movahed.m@northeastern.edu)

14 May 2022

---

# Table of contents

[Table of contents](#)

[Introduction](#)

[Analytics](#)

[Step 1: Split the data into a train and test set](#)

[Step 2: Estimating the lambda.min and lambda.1se values.](#)

[Step 3: Plot the results from the cv.glmnet function](#)

[Step 4: Fitting a Ridge regression model against the training set and reporting on the coefficients.](#)

[Step 5: Determine the performance of the fit model against the training set by calculating the root mean square error \(RMSE\)](#)

[Step 6: Determine the performance of the fit model against the test set by calculating the root mean square error \(RMSE\)](#)

[Step 7 and 8: estimating the lambda.min and lambda.1se values. And creating a plot for glmnet.](#)

[Step 9: Fit a LASSO regression model against the training set and report the coefficients.](#)

[Step 10: Determine the performance of the fit model against the training set by calculating the root mean square error \(RMSE\).](#)

[Step 11: Determine the performance of the fit model against the test set by calculating the root mean square error \(RMSE\)](#)

[Conclusion](#)

[Step 12: Which model performed better and why?](#)

[Step 13: Perform stepwise selection and then fit a model](#)

[Bibliography](#)

[Appendix](#)

---

## Introduction

This R practice mainly focuses on Regularization. During this assignment, I will use the `Collage` dataset to build regularization models using Ridge and Lasso, perform `stepwise selection`, and then fit a model in 14 steps.

---

## Analytics

This part will go through the steps described in the assignment and use R to predict `grade.rate`.

### Step 1: Split the data into a train and test set

this step is straightforward coding.

```
#split the data
set.seed(96)
trainIndex<-sample(x = nrow(dataset) , size = nrow(dataset)*0.7)
train <- dataset[trainIndex,]
test <- dataset[-trainIndex,]

train_x <- model.matrix(Grad.Rate~.,train)[,-1]
test_x <- model.matrix(Grad.Rate~.,test)[,-1]
```

```
train_y <- train$Grad.Rate
test_y <- test$Grad.Rate
```

The results are two datasets for modeling and testing separately.

## Step 2: Estimating the lambda.min and lambda.1se values.

I use `cv.glmnet()` to estimate `min` and `1se` of lambda in this part.

```
cv.lasso <- cv.glmnet(train_x, train_y, nfolds = 10)
#object : min perdition error
log(cv.lasso$lambda.min)
log(cv.lasso$lambda.1se)
```

The resulting plot is shown below.

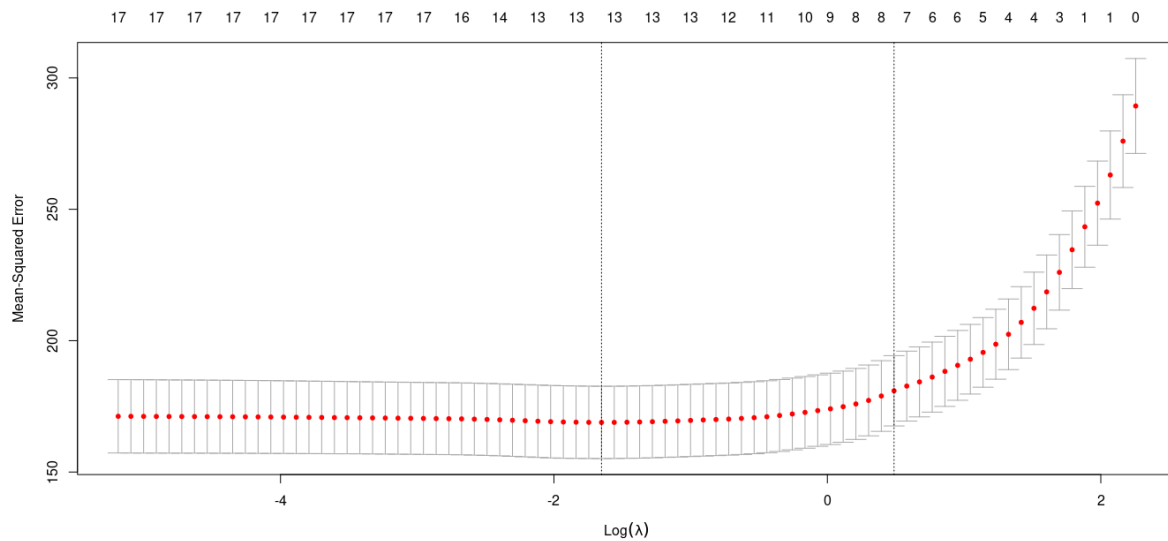
The resulting values for `lambda.min` and `lambda.1se` are shown below.

```
> #object : min perdition error
> log(cv.lasso$lambda.min)
[1] -1.652088
> log(cv.lasso$lambda.1se)
[1] 0.4876883
```

As it can be seen, the best model has 13 variables in it, while the simplest model that does the same thing has eight variables.

## Step 3: Plot the results from the cv.glmnet function

```
#finding the best value of lambda
set.seed(96)
cv.lasso <- cv.glmnet(train_x, train_y, nfolds = 10)
plot(cv.lasso)
```



the `cv.lasso` plot

As seen in the plot, we see a spike in the graph starts after setting lambda between 8 and 7 the mean of squared error starts to rise dramatically.

## Step 4: Fitting a Ridge regression model against the training set and reporting on the coefficients.

```
#fit model based on lambda.min with Ridge regression
model.1se <- glmnet(train_x,train_y,alpha = 0,lambda = cv.lasso$lambda.1se)
model.1se
#display reg coeff
coef(model.1se)
#display coeff of train model with no regularization
tm <- lm(Grad.Rate~.,data = train)
coef(tm)
```

The `coef(model.1se)` result is shown below.

```
> coef(model.1se)
18 x 1 sparse Matrix of class "dgCMatrix"
      s0
(Intercept) 3.209466e+01
PrivateYes  4.545804e+00
Apps        5.894800e-04
Accept      3.213770e-05
Enroll      3.249355e-04
Top10perc   8.051476e-02
Top25perc   1.127033e-01
F.Undergrad 5.744338e-05
P.Undergrad -1.254433e-03
Outstate    7.413867e-04
Room.Board  1.965067e-03
Books       -6.233951e-03
Personal    -2.017028e-03
PhD          6.396792e-02
Terminal    -8.927667e-03
S.F.Ratio   1.837197e-01
```

```
perc.alumni 2.659121e-01
Expend      -2.363661e-04
```

Compering these with coefficients of the full model shows a huge difference in Intercept; however, for most existing variables, the coefficients aren't that different between the two models

## Step 5: Determine the performance of the fit model against the training set by calculating the root mean square error (RMSE)

```
#train set prediction
preds.train<- predict(model.lse,newx = train_x)
train.rmse <- rmse(train_y,preds.train)
```

## Step 6: Determine the performance of the fit model against the test set by calculating the root mean square error (RMSE)

```
#test set prediction
preds.test <- predict(model.lse,newx =test_x)
test.rmse <- rmse(test_y,preds.test)
```

```
#view RMSE of full model
preds.tm <- predict(tm,new = test)
full.rmse <-rmse(test$Grad.Rate,preds.tm)
#compering values of rmse
train.rmse
test.rmse
full.rmse
```

The results are shown below

```
> train.rmse
[1] 12.6653
> test.rmse
[1] 12.79758
> full.rmse
[1] 12.8865
```

As can be seen, the regularized model performed better in reducing RMSE than the initial model.

## Step 7 and 8: estimating the lambda.min and lambda.1se values. And creating a plot for glmnet.

It is as same as steps 2 and 3; therefore, I skip it.

## Step 9: Fit a LASSO regression model against the training set and report the coefficients.

```
#fit model based on lambda.min with lasso regression
model.1se <- glmnet(train_x,train_y,alpha = 1,lambda = cv.lasso$lambda.1se)
model.1se
#display reg Coeff
coef(model.1se)
```

The results are shown below

```
> coef(model.1se)
18 x 1 sparse Matrix of class "dgCMatrix"
              s0
(Intercept)  4.001468e+01
PrivateYes   .
Apps         5.397571e-05
Accept       .
Enroll       .
Top10perc    3.591467e-02
Top25perc    1.128799e-01
F.Undergrad  .
P.Undergrad  -2.024551e-04
Outstate     1.020005e-03
Room.Board   9.719135e-04
Books        .
Personal     -1.011774e-03
PhD          .
Terminal     .
S.F.Ratio    .
perc.alumni  2.067628e-01
Expend       .
```

As seen in the results, nine variables from the original model lost all the coefficients in the new model and were removed.

## Step 10: Determine the performance of the fit model against the training set by calculating the root mean square error (RMSE).

```
#train set prediction
preds.train<- predict(model.1se,newx = train_x)
train.rmse <- rmse(train_y,preds.train)
```

## Step 11: Determine the performance of the fit model against the test set by calculating the root mean square error (RMSE)

```
#test set prediction
preds.test <- predict(model.1se,newx =test_x)
test.rmse <- rmse(test_y,preds.test)

#compering values of rmse
train.rmse
test.rmse
full.rmse
```

The results are shown below

```
> train.rmse
[1] 13.28165
> test.rmse
[1] 13.16788
> full.rmse
[1] 12.8865
```

As can be seen, the RMSE of the model increased; however, we managed to eliminate nine variables without making much RMSE in test data. The RMSE in testing data is still the same as training data, indicating the model's success.

## Conclusion

### Step 12: Which model performed better and why?

Each model has its unique advantages. The ridge model reduced RMSE more significantly; however, it used all the variables of the initial model; however, the lasso eliminated some variables making the model more simple.

### Step 13: Perform stepwise selection and then fit a model

```
# Forward selection method
tm <- lm(Grad.Rate~1., data = train)
datan <- (train[, unlist(lapply(train, is.numeric))])
m = cor(datan, datan$Grad.Rate)
#first rearranging cor matrix and remove duplicates
msort <- m %>%
  as.data.frame() %>%
  mutate(var1 = rownames(.)) %>%
  gather(var2, value, -var1) %>%
  arrange(desc(value)) %>%
  group_by(value) %>%
  filter(row_number()==1)

msort$value <- abs(msort$value)
msort <- msort[(msort$value < 0.999),]
msort <- arrange(msort, desc(value))
head(msort, 5)
steps <- step(tm, direction = 'forward', scope ~
  Outstate + perc.alumni + Top10perc + Top25perc + Room.Board)
model_forward <- lm(formula = Grad.Rate ~ Outstate + Top25perc + perc.alumni + Room.Board, data = train)
summary(model_forward)
#view RMSE of full model
preds.model_forward <- predict(model_forward, new = test)
full.rmse <- rmse(test$Grad.Rate, preds.model_forward)
full.rmse
```

the RMSE of model\_forward is 13.0137, which is very close to the ridge model, indicating the effectiveness of this method.

# Bibliography

Instructure.com. (2019). *Module 4 Assignment — Regularization*. [online] Available at: <https://northeastern.instructure.com/courses/116151/assignments/1369378> [Accessed 14 May 2022].

Pumpkin C (2022). *correlation of one variable to all the other in R*. [online] Stack Overflow. Available at: <https://stackoverflow.com/questions/45892274/correlation-of-one-variable-to-all-the-other-in-r> [Accessed 14 May 2022].

Guru99. (2020). *R Stepwise & Multiple Linear Regression [Step by Step Example]*. [online] Available at: <https://www.guru99.com/r-simple-multiple-linear-regression.html> [Accessed 14 May 2022].

---

## Appendix

```
install.packages('FSA')
install.packages('FSAdata')
install.packages('magrittr')
install.packages('dplyr')
install.packages('tidyr')
install.packages('plyr')
install.packages('tidyverse')
install.packages('outliers')
install.packages('ggplot2')
install.packages('lubridate')
install.packages('corrplot')

library(ggplot2)
library(outliers)
library(FSA)
library(FSAdata)
library(magrittr)
library(dplyr)
library(tidyr)
library(dplyr)
library(tidyverse)
library(scales)
library(lubridate)
library(corrplot)
install.packages('caret')
library(caret)
install.packages('ggplot2')
library(ggplot2)
install.packages("glmnet")
install.packages('Metrics')
library(glmnet)
library(Metrics)

#load the data
library(ISLR)
data("College")
dataset <- as.data.frame(College)
```



```

#split the data
set.seed(96)
trainIndex<-sample(x = nrow(dataset) , size = nrow(dataset)*0.7)
train <- dataset[trainIndex,]
test <- dataset[-trainIndex,]

train_x <- model.matrix(Grad.Rate~.,train)[-1]
test_x <- model.matrix(Grad.Rate~.,test)[-1]

train_y <- train$Grad.Rate
test_y <- test$Grad.Rate

#finding value of lambda

#finding the best value of lambda
set.seed(96)
cv.lasso <- cv.glmnet(train_x,train_y,nfolds = 10)
plot(cv.lasso)

#object : min perdition error
log(cv.lasso$lambda.min)
log(cv.lasso$lambda.1se)
#Ridge
#fit model based on lambda.min with Ridge regression
model.1se <- glmnet(train_x,train_y,alpha = 0,lambda = cv.lasso$lambda.1se)
model.1se
#display reg coeff
coef(model.1se)

#display coeff of train model with no regularization
tm <- lm(Grad.Rate~.,data = train)
coef(tm)

#view RMSE of full model
preds.tm <- predict(tm,new = test)
full.rmse <-rmse(test$Grad.Rate,preds.tm)

#train set prediction
preds.train<- predict(model.1se,newx = train_x)
train.rmse <- rmse(train_y,preds.train)
#test set prediction
preds.test <- predict(model.1se,newx =test_x)
test.rmse <- rmse(test_y,preds.test)

#compering values of rmse
train.rmse
test.rmse
full.rmse
#lasso
#fit model based on lambda.min with lasso regression
model.1se <- glmnet(train_x,train_y,alpha = 1,lambda = cv.lasso$lambda.1se)
model.1se
#display reg coeff
coef(model.1se)

#display coeff of train model with no regularization
tm <- lm(Grad.Rate~.,data = train)
coef(tm)

#view RMSE of full model
preds.tm <- predict(tm,new = test)
full.rmse <-rmse(test$Grad.Rate,preds.tm)

#train set prediction
preds.train<- predict(model.1se,newx = train_x)
train.rmse <- rmse(train_y,preds.train)
#test set prediction

```

```

preds.test <- predict(model.lse,newx =test_x)
test.rmse <- rmse(test_y,preds.test)

#compering values of rmse
train.rmse
test.rmse
full.rmse

# Forward selection method
tm <- lm(Grad.Rate~1.,data = train)
datan <- (train[, unlist(lapply(train, is.numeric))])
m = cor(datan,datan$Grad.Rate)
#first rearrenging cor matrix and remove duplicates
msort <- m %>%
  as.data.frame() %>%
  mutate(var1 = rownames(.)) %>%
  gather(var2, value, -var1) %>%
  arrange(desc(value)) %>%
  group_by(value) %>%
  filter(row_number()==1)

msort$value <- abs(msort$value)
msort<-msort[(msort$value < 0.999),]
msort <-arrange(msort,desc(value))
head(msort,5)
steps <- step(tm, direction = 'forward',scope ~
  Outstate + perc.alumni + Top10perc + Top25perc + Room.Board)
model_forward <- lm(formula = Grad.Rate ~ Outstate + Top25perc + perc.alumni + Room.Board, data = train)
summary(model_forward)
#view RMSE of full model
preds.model_forward <- predict(model_forward,new = test)
full.rmse <-rmse(test$Grad.Rate,preds.model_forward)
full.rmse

```